

```

#!/usr/bin/env python
# coding: utf-8

import os
# import args
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import random
from datetime import datetime, timedelta
from functools import reduce
# from sklearn.model_selection import train_test_split

def line_plot(df, var, title=False):
    fig, ax = plt.subplots(figsize=(10,5), dpi=100)
    plt.style.use('seaborn-ticks') # print(plt.style.available)
    if title:
        plt.title(title, fontsize=22)

    df.reset_index().plot.line(x= 'date', y = var, ax=ax, color="black").grid(alpha=0.5)
    ax.get_legend().remove()
#     fig.savefig(_charts + "\\\" + segname + "_MON.png", bbox_inches = 'tight', pad_inches=0)

def get_return(df, var):
    df['G1'] = df[var].pct_change(1)
    df['N1'] = np.log(df[var].pct_change(1)+1)

    df['G5'] = df[var].pct_change(5)
    df['N5'] = np.log(df[var].pct_change(5)+1)

    df['G1_mon'] = df[var].pct_change(25)
    df['N1_mon'] = np.log(df[var].pct_change(25)+1)

    df['G6_mon'] = df[var].pct_change(125)
    df['N6_mon'] = np.log(df[var].pct_change(125)+1)

    # df['G1_yr'] = df[var].pct_change(250)
    # df['N1_yr'] = np.log(df[var].pct_change(250)+1)

    # df['G2_yr'] = df[var].pct_change(500)
    # df['N2_yr'] = np.log(df[var].pct_change(500)+1)

#     df['MA2'] = df[var].rolling(2, min_periods=2).mean()
#     df['MA32'] = df[var].rolling(32, min_periods=1).mean()
#     df['MA16'] = df[var].rolling(16, min_periods=1).mean()
#     df['MA4'] = df[var].rolling(4, min_periods=1).mean()
#     df['EMA'] = df[var].ewm(alpha=0.5).mean()
    return df[126:]
#     return df

def get_return_std(df, y):
    df['R12_std'] = df[y].rolling(12,min_periods = 2).std()
    df['R25_std'] = df[y].rolling(25,min_periods = 2).std()

```

```

df['R50_std'] = df[y].rolling(50, min_periods = 2).std()
return df[2:]

```

```

def get_holidays(df):
    date = datetime.strptime(df.date.min(), "%Y-%m-%d")
    end_date = datetime.strptime(df.date.max(), "%Y-%m-%d")
    holidays = []
    while date <= end_date:
        if date.weekday() <= 4:
            date = date.strftime("%Y-%m-%d")
            if date not in list(df.date):
                holidays += [date]
            date = datetime.strptime(date, "%Y-%m-%d")
        date += timedelta(1)
    return holidays

```

--- Feature Engineering: return, s.d. of return, weekday, holiday, etc. ---

```

def get_state(df):
    df = get_return(df, 'close')
    df = get_return_std(df, 'N1')
    # keep states needed
    df = df[['date', 'close', 'vol'] + list(df.filter(regex='^N|^R|^MA'))]
    return df

```

```

def get_training_set(dic):
    train = {}
    for key, value in dic.items():
        # train[key] = value.loc[(value.date >= start_date) & (value.date < end_date)]
        train[key] = value
        train[key] = get_state(train[key])
        train[key] = train[key].iloc[:2200].reset_index(drop=True)
        train[key]['weekday'] = pd.to_datetime(train[key].date).dt.dayofweek
    return train

```

save 500 days to calculate 2-year return

```

def get_test_set(dic):
    test = {}
    for key, value in dic.items():
        # test[key] = value.iloc[int(value.index[value.date==end_date].values) - 500:]
        test[key] = value
        test[key] = get_state(test[key])
        test[key] = test[key].iloc[2200:].reset_index(drop=True)
        test[key]['weekday'] = pd.to_datetime(test[key].date).dt.dayofweek
    return test

```

```

def get_holiday_dummy(dic):
    one_day_before_holiday = list((pd.to_datetime(holidays) - timedelta(1)).strftime("%Y-%m-%d"))
    two_day_before_holiday = list((pd.to_datetime(holidays) - timedelta(2)).strftime("%Y-%m-%d"))
    three_day_before_holiday = list((pd.to_datetime(holidays) - timedelta(3)).strftime("%Y-%m-%d"))
    three_day_after_holiday = list((pd.to_datetime(holidays) + timedelta(3)).strftime("%Y-%m-%d"))
    two_day_after_holiday = list((pd.to_datetime(holidays) + timedelta(2)).strftime("%Y-%m-%d"))
    one_day_after_holiday = list((pd.to_datetime(holidays) + timedelta(1)).strftime("%Y-%m-%d"))

    for key, value in dic.items():
        df = dic[key]

```

```

df['before_holiday_1'] = df.date.apply(lambda x: 1 if (x in one_day_before) else 0)
df['before_holiday_2'] = df.date.apply(lambda x: 1 if (x in two_day_before) else 0)
df['before_holiday_3'] = df.date.apply(lambda x: 1 if (x in three_day_before) else 0)
df['after_holiday_1'] = df.date.apply(lambda x: 1 if (x in one_day_after_holiday) else 0)
df['after_holiday_2'] = df.date.apply(lambda x: 1 if (x in two_day_after_holiday) else 0)
df['after_holiday_3'] = df.date.apply(lambda x: 1 if (x in three_day_after_holiday) else 0)
return dic

```

```

# load stocks
# data_folder = args._daily
data_folder = '/content/stock_prediction/data/nifty'
os.chdir(data_folder)
stock = [x.replace('.csv', '') for x in os.listdir(data_folder)]
# stock = [x for x in stock if x!='GOOGL' and x!='ALLY' and x!='WAL']

# stock = ['SPX', 'BA', 'MS']
tables = [pd.read_csv(x + ".csv") for x in stock if x!='.ipynb_checkpoints']
os.chdir('/content')
# stock = ['SBIN', 'nifty']
# tables = [pd.read_csv(x + ".csv") for x in stock ]
# dt_sbi = pd.read_csv('SBIN.csv')
# dt_sbi = dt_sbi[:3105]
# dt_nifty = pd.read_csv('nifty.csv')

# df_nifty = pd.DataFrame(columns=dt_nifty.columns)
# df_sbi = pd.DataFrame(columns=dt_nifty.columns)

# for i in range(len(dt_sbi)):
#     df_nifty.loc[i] = dt_nifty.loc[len(dt_nifty)-i-1]
#     df_sbi.loc[i] = dt_sbi.loc[len(dt_sbi)-i-1]

# tables = [df_sbi, df_nifty]
# tables = pd.read_csv(stock+'.csv')
stock_d = dict(zip(stock, tables)) # line_plot(stock_d['SPX'], 'close', 'SPX')
# clean a bit
# print(stock_d['SBIN'])
for each in stock:
    # stock_d[each] = stock_d[each].rename(columns={'4. close': 'close', '5. volume': 'volume'})
    stock_d[each] = stock_d[each].rename(columns={'Date': 'date', 'Adj Close': 'close'})

# env para
start_date = max([value.date.min() for key, value in stock_d.items()])
end_date = '2020-04-21'
worth = 100000
bottom_line = worth * 0.95
cycle = 60

# get train and test
train = get_training_set(stock_d)
test = get_test_set(stock_d)
# holidays = get_holidays(train['SPX']) + get_holidays(test['SPX'])
# train = get_holiday_dummy(train)
# test = get_holiday_dummy(test)
# train['SPX']['close'] = train['SPX']['close']/10

```

```

for key, df in train.items():
    train[key].fillna(method = 'backfill',inplace = True)

for key, df in test.items():
    test[key].fillna(method = 'backfill',inplace = True)


def add_days(date, days):
    new_datetime = datetime.strptime(date, "%Y-%m-%d") + timedelta(days)
    new_date = new_datetime.strftime("%Y-%m-%d")
    return new_date

def add_one_trade_day(date):
    date = add_days(date, 1)
    # while (datetime.strptime(date, "%Y-%m-%d").weekday()>=5) or (date in holiday:
    while (datetime.strptime(date, "%Y-%m-%d").weekday()>=5):
        date = add_days(date, 1)
    return date

def concat(x,y): return pd.concat([x,y], axis=0).reset_index(drop=True)

def get_info(env, date):
    info_matrix = []
    for key, df in env.items():
        df['stock'] = key
        a = df.loc[df.date == date]
        info_matrix += [a]
    return reduce((lambda x, y: concat(x,y)), info_matrix)

def fraction_power(x, n):
    if 0<=x: return x**(1./n)
    return -(-x)**(1./n)

# this env does not have interest rate or mutual fund, and has no transaction cost
class trade_env():
    def __init__(self, env, worth, cycle, rho):
        self.env = env
        self.cycle = cycle
        self.rho = rho
        self.action_space = len(env.keys()) + 1 # add one to add cash holding
        # self.state_space = len(env['nifty'].columns) - 1 # minus one to remove d
        # self.state_space = len(env['SBIN'].columns) - 1
        self.state_space = len(env['^NSEI'].columns) - 1
        self.end_date = env['^NSEI'].date.max()
        # self.end_date = '2020-04-19'
        self.reset(worth)

    def step(self, action):

        # today
        last_worth = self.worth
        last bse worth = self.bse worth

```

```

info = get_info(self.env, self.date)

# SPX_price = info.loc[info.stock == 'SPX'].close.item()
# SPX_share = self.SPX_worth // SPX_price
# SPX_changes = self.SPX_worth % SPX_price

# nifty_price = info.loc[info.stock == 'nifty'].close.item()
# if len(info[info.stock == 'nifty'].close)==1:
#     nifty_price = info[info.stock=='nifty'].close.item()
# else:
#     nifty_price = last

i1 = 0
while(len(info[info.stock=='^NSEI'].close)!=1 and self.date<=end_date):
    i1 = i1+1
    if(i1>1000):
        print('step function first loop running too many times\n')
        break
    self.date = add_one_trade_day(self.date)
    info = get_info(self.env, self.date)

bse_price = info[info.stock == '^NSEI'].close.item()
# nifty_price = info[info.stock == 'nifty'].close.item()
# nifty_price = info[info.stock=='nifty'].close[0]
bse_share = self.bse_worth // bse_price
bse_changes = self.bse_worth % bse_price

self.prices = np.append([1], info.close.to_numpy()) # 1 is the opportunity cost
self.share = self.worth * action // self.prices
changes = (self.worth * action % self.prices)[1: self.action_space]
self.share[0] = self.share[0] + np.sum(changes) # changes go to cash, i.e.

# next day

self.date = add_one_trade_day(self.date)
new_info = get_info(self.env, self.date)

i2 = 0
while(len(new_info[new_info.stock=='^NSEI'].close)!=1 and self.date<=end_date):
    i2 = i2 + 1
    if(i2>1000):
        print("step function second loop running too many times\n")
        break
    self.date = add_one_trade_day(self.date)
    new_info = get_info(self.env, self.date)

bse_new_price = new_info[new_info.stock == '^NSEI'].close.item()

# nifty_new_price = new_info.loc[new_info.stock == 'nifty'].close.item()
# nifty_new_price = new_info[new_info.stock == 'nifty'].close.item()
# nifty_new_price = new_info[new_info.stock == 'nifty'].close[0]
# if len(new_info[new_info.stock == 'nifty'].close)>1:

```

```

# nifty_new_price = new_info[new_info.stock=='nifty'].close[0]
# elif len(new_info[new_info.stock == 'nifty'].close)==1:
# nifty_new_price = new_info[new_info.stock == 'nifty'].close.item()
# else:
# print("GOGHEGHEHGOEHIGOHGNLKSNNVLNVNL\n")
self.bse_worth = bse_share * bse_new_price + bse_changes
bse_reward = self.bse_worth - last_bse_worth

new_prices = np.append([1], new_info.close.to_numpy())
self.worth = np.sum(self.share * new_prices)

# reality
state = np.array(new_info.drop(['date', 'stock'],axis=1))
reward = self.worth - last_worth
utility = fraction_power(reward, self.rho)
gone = datetime.strptime(self.date, "%Y-%m-%d") - datetime.strptime(self.i

if (self.worth < bottom_line) or (gone.days > self.cycle) or (self.date ==
    done = True
else:
    done = False

return state, reward, bse_reward, utility, done

def reset(self, worth):
    self.worth = worth
    self.bse_worth = worth
    self.date = random.choice(self.env['^NSEI'].date)
    self.initial_date = self.date
    initial_info = get_info(self.env, self.initial_date)
    # i=0
    while(len(initial_info.close)==0):
        i=i+1
        if(i>1000):
            print("Reset running too many times\n")
            # break
        self.date = random.choice(self.env['^NSEI'].date)
        self.initial_date = self.date
        initial_info = get_info(self.env, self.initial_date)
    # print(initial_info)
    initial_state = np.array(initial_info.drop(['date', 'stock'],axis=1))

    return initial_state

```



EVERYTHING ABOVE THIS WAS ENV.PY

```
from numpy.random import randn
import numpy as np
from itertools import count
from collections import namedtuple
from torch.utils.tensorboard import SummaryWriter
import os
import seaborn as sns
# import args
from scipy import stats
import pandas as pd

from replay_memory import Memory, Transition
from ounoise import OrnsteinUhlenbeckActionNoise as noise
# from env import trade_env
# import env

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

import matplotlib.pyplot as plt
%matplotlib inline

address = '/content/stock_prediction'
_raw = address + '/raw'
_daily = _raw + '/daily'
_weekly = _raw + '/weekly'

# address = args.address
writer = SummaryWriter(address + '/tensor')
weights = address + '/weights'
tensors = address + '/tensor'
outputs = address + '/outputs'

seed = 543
memory = Memory(100000)
noise_scale = 1.5
final_noise_scale = 0.5
worth = 1000000
cycle = 7
```

```
rho = 3 # coefficient of risk aversion
```

```
# for file in os.listdir(tensors):  
#     file = os.path.join(tensors,file)  
#     os.remove(file)
```

```
class Actor(nn.Module):
```

```
    def __init__(self, hidden_size, num_inputs, action_space):  
        super(Actor, self).__init__()  
        self.affine1 = nn.Linear(num_inputs, hidden_size)  
        self.ln1 = nn.LayerNorm(hidden_size)
```

```
        self.affine2 = nn.Linear(hidden_size, hidden_size)  
        self.ln2 = nn.LayerNorm(hidden_size)
```

```
        self.value = nn.Linear(hidden_size, action_space)
```

```
    def forward(self, x):
```

```
        x = self.affine1(x)  
        x = self.ln1(x)  
        x = F.relu(x)
```

```
        x = self.affine2(x)  
        x = self.ln2(x)  
        actions = F.softmax(self.value(x), dim=-1)
```

```
        return actions
```

```
class Critic(nn.Module):
```

```
    def __init__(self, hidden_size, num_inputs, action_space):  
        super(Critic, self).__init__()  
        self.affine1 = nn.Linear(num_inputs, hidden_size)  
        self.ln1 = nn.LayerNorm(hidden_size)
```

```
        self.affine2 = nn.Linear(action_space + hidden_size, hidden_size)  
        self.ln2 = nn.LayerNorm(hidden_size)
```

```
        self.value = nn.Linear(hidden_size, 1)
```

```
    def forward(self, x, actions):
```

```
        x = self.affine1(x)  
        x = self.ln1(x)  
        x = F.relu(x)
```

```
        x = torch.cat((x, actions), 1)  
        x = self.affine2(x)  
        x = self.ln2(x)  
        q_value = self.value(F.relu(x))
```

```
        return q_value
```

```
# torch.manual_seed(seed)
```



```

# train_env = trade_env(env.train, worth, cycle, rho)
# test_env = trade_env(env.test, worth, cycle, rho)

train_env = trade_env(train, worth, cycle, rho)
test_env = trade_env(test, worth, cycle, rho)

hidden_size = 128
action_space = train_env.action_space
state_space = train_env.state_space
num_inputs = (action_space - 1) * state_space # flatten info matrix; minus one si

critic = Critic(hidden_size, num_inputs, action_space)
critic_target = Critic(hidden_size, num_inputs, action_space)
actor = Actor(hidden_size, num_inputs, action_space)
actor_target = Actor(hidden_size, num_inputs, action_space)
actor_perturbed = Actor(hidden_size, num_inputs, action_space)

critic_optim = optim.Adam(critic.parameters(), lr=1e-3)
actor_optim = optim.Adam(actor.parameters(), lr=1e-3)

eps = np.finfo(np.float32).eps.item()

noise = noise(action_space)
# noise.reset()

tau = 0.001
def soft_update(target, source, tau):
    for target_param, param in zip(target.parameters(), source.parameters()):
        target_param.data.copy_(target_param.data * (1.0 - tau) + param.data * tau)
def hard_update(target, source):
    for target_param, param in zip(target.parameters(), source.parameters()):
        target_param.data.copy_(param.data)

def select_action_without_noise(state):
    action = actor(state)
    action = action.detach().numpy()
    return action

# select action with para noise on the last layer
def select_action(state):

    hard_update(actor_perturbed, actor)
    actor_params = actor_perturbed.state_dict()

    param = actor_params['value.bias']
    param += torch.tensor(noise()).float()

    action = actor_perturbed(state)
    return action.detach().numpy()

def select_action_with_para_noise(state):

```

```

hard_update(actor_perturbed, actor)
actor_params = actor_perturbed.state_dict()

param = actor_params['value.bias']
param += torch.tensor(randn(action_space) * 1.8).float()

action = actor_perturbed(state)
return action.detach().numpy()

```

```

hard_update(critic_target, critic)
gamma = 0.99

```

```

def update_para():
    transitions = memory.sample(128)
    batch = Transition(*zip(*transitions))

    state_batch = torch.stack(batch.state)
    action_batch = torch.stack(batch.action)
    utility_batch = torch.stack(batch.utility)
    mask_batch = torch.stack(batch.mask)
    next_state_batch = torch.stack(batch.next_state)

    q_batch = critic(state_batch, action_batch)
    next_action_batch = actor_target(next_state_batch)
    next_q_batch = utility_batch + gamma * mask_batch * critic_target(next_state_batch)

    value_loss = F.mse_loss(q_batch, next_q_batch)
    critic_optim.zero_grad()
    value_loss.backward()
    critic_optim.step()

    policy_loss = - critic(state_batch, actor(state_batch))
    policy_loss = policy_loss.mean()
    actor_optim.zero_grad()
    policy_loss.backward()
    actor_optim.step()

    soft_update(actor_target, actor, tau)
    soft_update(critic_target, critic, tau)
    return value_loss, policy_loss

```

```

def save_model(path):
    if not os.path.exists(path):
        os.makedirs(path)

    ddp_actor = os.path.join(path, 'ddpg_actor_weights')
    ddp_critic = os.path.join(path, 'ddpg_critic_weights')
    torch.save(actor.state_dict(), ddp_actor)
    torch.save(critic.state_dict(), ddp_critic)

```

```

def get_state_tensor(state):
    state = state.reshape(-1, num_inputs)[0]

```

```

return torch.tensor(state).float()

updates = 0
running_reward = []

# for i_episode in range(5000):
for i_episode in range(1000):
    print("Episode "+str(i_episode)+" running")
    abs_training_reward = 0
    relative_training_reward = 0
    state = get_state_tensor(train_env.reset(worth))

    # dwindling noise
    noise.scale = (noise_scale - final_noise_scale) * max(0, 3000-i_episode)/3000
    if train_env.date!=train_env.end_date:
        for t in range(1, 400):

            action = select_action(state)
            # print(action,'\n')
            next_state, reward, bse_reward, utility, done = train_env.step(action)
            # print(bse_reward,'\n')
            # print(action)
            # print(train_env.worth)

            abs_training_reward += reward
            relative_training_reward += (reward - bse_reward)
            action = torch.tensor(action).float()
            mask = torch.tensor([not done]).float()
            reward = torch.tensor([reward])
            utility = torch.tensor([utility])
            next_state = get_state_tensor(next_state)

            # save to memory
            memory.push(state, action, mask, utility, next_state)

            if len(memory)>=128:
                for _ in range(4):
                    value_loss, policy_loss = update_para()
                    writer.add_scalar('loss/value', value_loss, updates)
                    writer.add_scalar('loss/policy', policy_loss, updates)
                    updates += 1

            state = next_state

            if done:
                break
            # print(abs_training_reward)
            writer.add_scalar('training/abs reward', abs_training_reward, i_episode)
            writer.add_scalar('training/relative reward', relative_training_reward, i_

# test sample, evaluate model performance
abs_test_reward = 0
bse_test_reward = 0
if i_episode % 1 == 0:

```

```

state = get_state_tensor(test_env.reset(worth))
if test_env.date!=test_env.end_date:
    for t in range(1, 100):

        action = select_action_without_noise(state)
        # print(action,'\n')

        next_state, reward, bse_reward, utility, done = test_env.step(action)
        abs_test_reward += reward
        bse_test_reward += bse_reward

        next_state = get_state_tensor(next_state)
        state = next_state

    if done:
        break

    relative_test_reward = abs_test_reward - bse_test_reward
    writer.add_scalar('test/abs reward', abs_test_reward, i_episode)
    writer.add_scalar('test/relative reward', relative_test_reward, i_episode)

    running_reward += [relative_test_reward]
# print(running_reward,'\n')
if len(running_reward)>1000 and np.median(running_reward[-100:]) > 500: # i

    save_model(weights)
    break

```



bold text--- Test Models ---

```
# actor_weights = os.path.join(weights, 'ddpg_actor_weights')
# critic_weights = os.path.join(weights, 'ddpg_critic_weights')

# actor.load_state_dict(torch.load(actor_weights))
# critic.load_state_dict(torch.load(critic_weights))

reward_list = []
relative_reward_list = []
for i_episode in range(1000):

    print("Episode "+str(i_episode)+" running")
    state = get_state_tensor(test_env.reset(worth))
    episode_reward = 0
    relative_episode_reward = 0
    if test_env.date!=test_env.end_date:
        for t in range(1, 500):

            action = select_action_without_noise(state)
            #         print(test_env.share)
            #         print(test_env.date)

            next_state, reward, bse_reward, utility, done = test_env.step(action) ;
            # print(bse_reward,'\n')
            episode_reward += reward
            relative_episode_reward += (reward - bse_reward)
```

```
next_state = get_state_tensor(next_state)
state = next_state

if done:
    break
print("Episode {} reward: {}".format(str(test_env.date), str(episode_reward)))

reward_list += [episode_reward]
relative_reward_list += [relative_episode_reward]
```



```
df = pd.DataFrame({'Relative Reward':relative_reward_list,'Absolute Reward': reward_list})
df.to_csv(outputs + '/test_sample_return.csv', index=False)
```

```
ax = sns.distplot(relative_reward_list)
ax.set(title='Distribution of Relative Rewards',xlabel='Relative Rewards', ylabel='Density')
```



```
ax= plt.plot(relative_reward_list)
plt.xlabel('Episode Number')
plt.ylabel('Relative Reward')
plt.savefig('abc.jpg')
# ax.set(title='Convergence of Relative Rewards',xlabel='Episode Number', ylabel='Relative Reward')
```



```
np.nanmedian(relative_reward_list)
```



```
np.nanmean(relative_reward_list)
```



```
ax= sns.distplot(reward_list)  
ax.set(title='Distribution of Rewards',xlabel='Rewards', ylabel='Probability')
```



```
plt.plot(reward_list)  
plt.xlabel('Episode Number')  
plt.ylabel('Reward (Return)')  
plt.savefig('abc.jpg')
```



```
np.nanmedian(reward_list)
```



```
np.nanmedian(reward_list)
```



```
np.nanmean(reward_list)
```



```
np.nanstd(reward_list)
```



```
np.nanstd(relative_reward_list)
```



```
stats.ttest_1samp(np.array(relative_reward_list), 0)
```



```
bse_reward_list = np.array(reward_list) - np.array(relative_reward_list)
```

```
sns.distplot(bse_reward_list)
```



```
np.nanmedian(bse_reward_list)
```



```
np.nanmean(bse_reward_list)
```



```
np.nanstd(bse_reward_list)
```



▼ Train Env

```
# actor_weights = os.path.join(weights, 'ddpg_actor_weights')
# critic_weights = os.path.join(weights, 'ddpg_critic_weights')

# actor.load_state_dict(torch.load(actor_weights))
# critic.load_state_dict(torch.load(critic_weights))

train_reward_list = []
train_relative_reward_list = []
for i_episode in range(1000):
    print('Episode ' + str(i_episode) + ' running\n')
    state = get_state_tensor(train_env.reset(worth))
    episode_reward = 0
    relative_episode_reward = 0
    if train_env.date != train_env.end_date:
        for t in range(1, 100):

            action = select_action(state)
            # print(train_env.share)
            # print(train_env.bse_worth)

            next_state, reward, bse_reward, utility, done = train_env.step(action)
            episode_reward += reward
            relative_episode_reward += (reward - bse_reward)

            next_state = get_state_tensor(next_state)
            state = next_state

            if done:
                break
    # print("Episode {} reward: {}".format(str(train_env.date), str(episode_reward)))

    train_reward_list += [episode_reward]
    train_relative_reward_list += [relative_episode_reward]
```



```
sns.distplot(train_relative_reward_list)
```



```
np.median(train_relative_reward_list)
```



```
np.mean(train_relative_reward_list)
```



```
stats.ttest_1samp(np.array(reward_list), 0)
```



```
ax= sns.distplot(train_reward_list)  
ax.set(title='Distribution of Relative Rewards',xlabel='Relative Rewards', ylabel=
```



```
plt.plot(train_reward_list)  
plt.xlabel('Episode Number')  
plt.ylabel('Reward (Return)')  
plt.savefig('abc.jpg')
```



```
np.median(train_reward_list)
```



```
np.mean(train_reward_list)
```



```
np.std(train_reward_list)
```

